Sovereign AI Usage Policy

Document ID: APEL-AI-104 | Owner: AI Governance | Revision: 2.3

This page contains the policy-specific controls for allowed use.

Allowed use

- Internal summarization
- Knowledge retrieval
- Draft approval notes
- Code assistance in sandbox
- Spreadsheet analysis
- Multimodal document review

Data boundary

This page contains the policy-specific controls for data boundary.

Data boundary

- Confidential information stays inside approved infrastructure.
- External AI APIs are not used for confidential workflows.
- Local model inference and local retrieval are approved platform capabilities.

Agent permissions

This page contains the policy-specific controls for agent permissions.

Agent permissions

- Inspection Analyst may use OCR, vision, retrieval, calculation and document-generation tools.
- Coding Assistant may use source files and an isolated sandbox.
- Agents cannot self-approve high-risk actions.

Human oversight

This page contains the policy-specific controls for human oversight.

Human oversight

- AI output is labeled as generated content.
- High-risk recommendations enter an approval queue.
- A human reviewer owns the final decision.

Audit requirements

This page contains the policy-specific controls for audit requirements.

Audit requirements

- Task creation
- Model selection
- Retrieval
- Tool calls
- Artifact generation
- Approval decision are retained as auditable events.